

Runguo Li

runguo.ai@gmail.com

runguoli.com

github.com/LiRunGuo

Champaign (USA)

Research Interests

ML Systems for large-scale models: inference engines and expert streaming for MoE models, distributed training and rollout infrastructure (ZeRO-3, FSDP, Hybrid Engine), RL post-training systems, and GPU compiler and kernel-level correctness. I work between the machine and the model — how weights are read from storage, how collectives and thread teams are sized, and whether a resource plan and the runtime agree about what will actually be built — while keeping a research background in LLM reasoning, multimodal learning, and agent safety.

Education

University of Illinois Urbana-Champaign (UIUC)

M.S. in Information Science

Aug 2026–Jul 2028

Urbana-Champaign, IL, USA

Shanghai University of Finance and Economics (SUFU)

B.S. in Business Analytics

Sep 2022–Jun 2026

Shanghai, China

Open-Source Contributions (ML Systems)

Merged fixes to a pure-C MoE inference engine, DeepSpeed, Triton, verl, LMCache, vLLM-Omni, Apache TVM, and FlashAttention, plus open pull requests across the wider AI infrastructure stack.

colibri — pure-C MoE inference engine

[Merged pull requests](#)

multi-drive expert streaming for the 510 GB DeepSeek-V4.1 container, physical-core OpenMP team sizing in four engines that were 18.7× slower without it, and a plan/runtime mismatch that silently allocated 6.25 GiB of unbudgeted KV cache

DeepSpeed — distributed training

[Merged pull requests](#)

fixed a ZeRO-3 rollout deadlock caused by unsynchronized generation stopping, and blocked partial Hybrid Engine policy injection for unsupported architectures (validated on H200 and MI250 GPUs)

Triton — GPU compiler

[Merged pull requests](#)

stopped nested-loop fusion from trusting an `llvm.assume` outside the loop, which removed a zero-trip guard and could cause incorrect memory writes (MLIR regression test and H200 reproducer)

verl — RL post-training

[Merged pull requests](#)

restored FSDP value-head critic loading after TRL relocated its value-head model classes

LMCache — KV cache management

[Merged pull requests](#)

rejected non-positive allocations that corrupted the free list, and enforced lazy %-format logging (ruff G004) so new f-string logging can no longer land in already-migrated files

vLLM-Omni — multimodal serving

[Merged pull requests](#)

rejected a stray audio placeholder that killed the Qwen2.5-Omni thinker engine when audio is read from video, turning an engine crash into an upfront validation error

Apache TVM — deep-learning compiler

[Merged pull requests](#)

fixed the ONNX Hardmax lowering when the reduced axis has a symbolic extent, where `one_hot` requires a static depth but the importer passed a symbolic variable

FlashAttention — attention kernels

[Merged pull requests](#)

fit the SM90 CuTe forward tile to the block-sparse block size so every sparse layout is accepted, and skipped the causal mask on unmasked KV blocks, cutting the unmasked inner loop from 580 to 477 instructions on H200

Open pull requests. [PyTorch #196787](#), [vLLM #56198](#), [SGLang #39315](#), [Megatron-LM #7287](#), [TensorRT-LLM #18904](#), [CUTLASS #3613](#), [DeepSpeed #8390](#), [DeepSpeed #8392](#), [Miles #3266](#), [Miles #3267](#), [Miles #3268](#), [Miles #3269](#), [Cordis #148](#), [Open Code Review #1351](#).

[Full list](#)

Publications

(*Co-first author)

1. Zhi Yang*, **Runguo Li***, Qiqi Qiang, Jiashun Wang, Fangqi Lou, Mengping Li, Dongpo Cheng, et al. “**FinVault: Benchmarking Financial Agent Safety in Execution-Grounded Environments**”. *arXiv preprint arXiv:2601.07853*, 2026.

The first execution-grounded safety benchmark for financial agents: 31 regulatory sandbox scenarios with writable state, 107 real-world vulnerabilities, and 963 test cases. Current defenses do not transfer — attack success rates reach 50.0% and stay non-negligible (6.7%) even for the most robust model.

[PDF](#) | [Project page](#)

2. **Runguo Li*** et al. “**Reasoning-Visual Critical Token Fine-Tuning for Multimodal Reasoning**”. Manuscript under review at AAAI 2027.

RVCFIT selects which chain-of-thought tokens receive direct supervision using reasoning-relevance and visual-sensitivity signals while retaining the full response as context; at 50% token retention it gives the best average among the evaluated multimodal reasoning benchmarks.

[PDF](#) | [Project page](#)

3. **Runguo Li*** et al. “**VeriBRT: Plan-Guided, Evidence-Based Automated Bug Reproduction**”. Manuscript under review at ICSE 2027.

Represents issue-report behavior as a persistent Two-Axis Plan (a Checklist for where a test must reach, an Intentlist for what it must verify) with reliability-aware evidence and localization coverage as a validity criterion; reproduces 341 of 433 issues on SWT-Bench Verified, 9.7 points above the strongest same-model baseline.

[PDF](#) | [Project page](#)

Research Experience

University of Illinois Urbana-Champaign (UIUC)

Aug 2026–Present

Research Intern, advised by Prof. Minjia Zhang

Urbana-Champaign, IL, USA

- Research on efficient machine learning systems for large models: training and inference efficiency, model compression and sparsity, kernel- and hardware-aware optimization, and system support for LLM agents.
- Bring an upstream-first practice to this work: reproduce failures in production inference and training stacks, then contribute fixes with regression tests and measured before/after numbers.

Shanghai Jiao Tong University

Mar 2026–Jul 2026

Research Intern, advised by Prof. Xiaodong Gu

Shanghai, China

- Researched LLMs for code, with an emphasis on repository-level program understanding, automated bug reproduction, and coding-agent planning and execution.
- Developed **VeriBRT**, a plan-guided framework that represents issue-report behavior with a persistent Two-Axis Plan and uses reliability-aware evidence for bug-reproducing test generation; set up the evaluation harness on SWT-Bench Verified.

SUFE FinAI Center

Jul 2025–Jul 2026

Research Intern, advised by Prof. Liwen Zhang

Shanghai, China

- Co-developed **FinVault**, an execution-grounded benchmark for financial-agent safety with 31 regulatory scenarios, 107 real-world vulnerabilities, and 963 test cases.
- Built the executable sandboxes behind the benchmark — writable databases, permission and quota constraints, toolchains, and audit logs — so safety failures are judged from verified state transitions rather than generated text.
- Studied financial reasoning and agent safety through chain-of-thought data curation, supervised fine-tuning, GRPO training, and evaluation on FinEval, FinQA, and ConvFinQA.

Tencent Youtu AI Lab

Oct 2025–May 2026

Research Intern

Shanghai, China

- Researched LLM-driven adversarial data synthesis and low-resource, multi-label content-safety evaluation, and contributed to a unified multilingual moderation model on an XLM-RoBERTa backbone.
- Developed **RVCFT**, a selective-supervision method that identifies reasoning-relevant and visually sensitive tokens for multimodal chain-of-thought fine-tuning.

Selected Open-Source Project

ARH — AI Research Helper

github.com/LiRunGuo/Arhelper (MIT, 0.2.0-alpha)

- A CLI-first research assistant runtime: tool calling, plan/execute safety with approval-gated writes, three-layer memory (preference extraction, BM25 cross-session recall, reflection summaries), skill self-learning, checkpoint/resume, and automatic context compression.
- Multi-LLM routing with cooldown-based failover across OpenAI, Anthropic, and Ollama; Hub-Spoke gateways (CLI, Telegram, FastAPI); research-native tools for arXiv, HuggingFace, GitHub, and OpenReview. Python 3.9+, Pydantic, FastAPI, SQLAlchemy; all state local under `~/ .arh/`.

Technical Skills

- **ML Systems and Inference:** MoE expert streaming and disk-resident inference, KV-cache and memory budgeting, attention and quantization kernels, serving engines (vLLM, SGLang, TensorRT-LLM), CUDA/Triton kernel-level debugging, GPU compiler passes (MLIR), and performance profiling.
- **Training and Post-Training Infrastructure:** distributed data/tensor/pipeline parallelism, ZeRO-3 and FSDP, Megatron-LM, DeepSpeed Hybrid Engine, collective synchronization and deadlock debugging, RL rollout systems (verl, TRL), SFT and preference optimization (DPO/GRPO/PPO), distillation, and LoRA/PEFT.
- **Languages and Tooling:** Python, C, CUDA/Triton, SQL, Bash, L^AT_EX; PyTorch, Transformers, FlashAttention, FAISS; Git-based upstream contribution, regression testing, reproducible benchmarking, Linux, and Docker.
- **Systems for Agents:** agentic planning and tool use, RAG and long-term memory, execution-grounded evaluation, runtime safety gates, and multi-channel gateways (FastAPI).

Honors & Awards

National Encouragement Scholarship

|

People’s Scholarship (2nd Class)

|

Silver Medal, FLTRP English Competition (Municipal)